



Marilyn Waldman

@mdwaldman

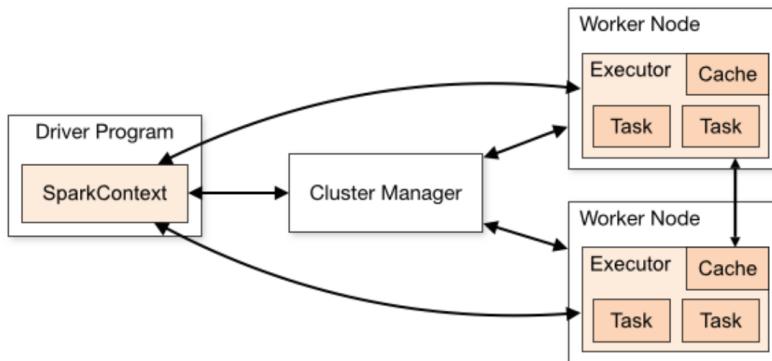
Agenda

- **What is Spark - Why Spark**
- **lab 1 - Functional programming -map and reduce**
- **lab 2 - RDD's**
- **lab 3 - Word Count**
- **lab 4 - Spark SQL**
- **Conclusions**

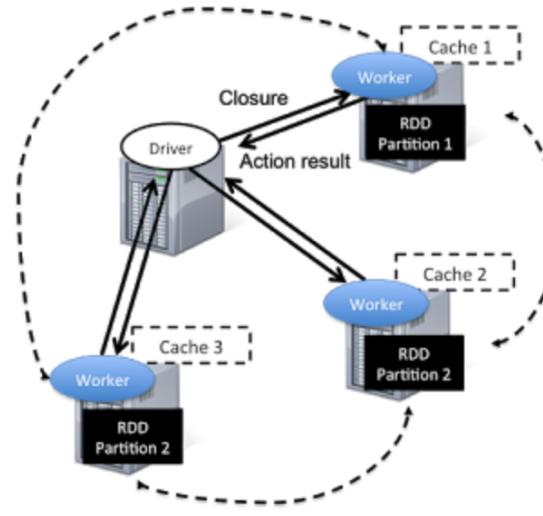




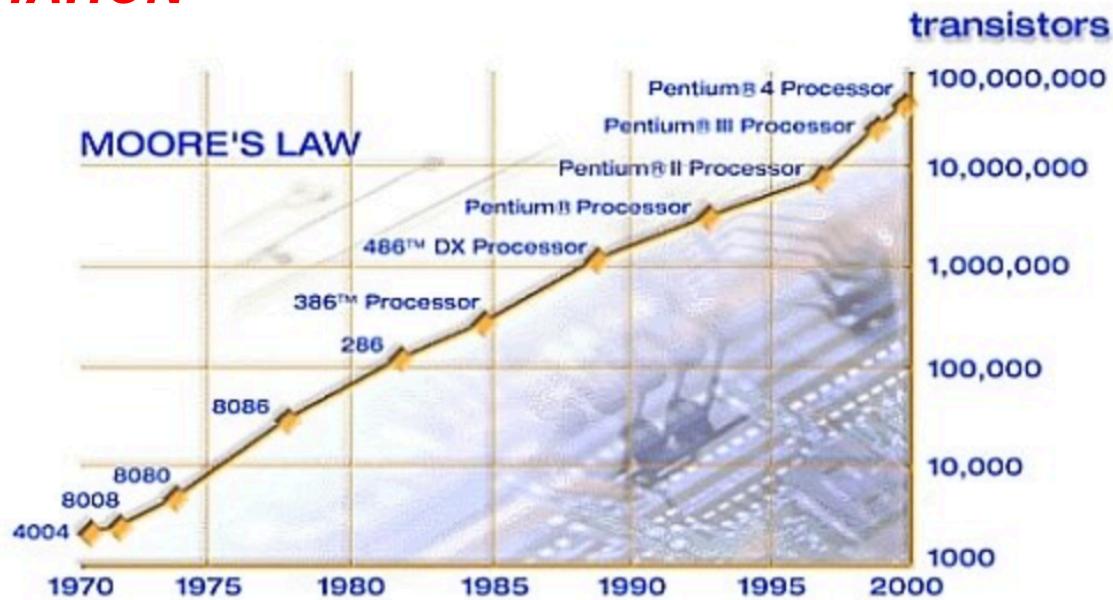
Cluster computing platform, a *distributed system*



- genomics
- regression - optimization
- real-time data
- anomaly detection
- fraud detection
- codeneuro

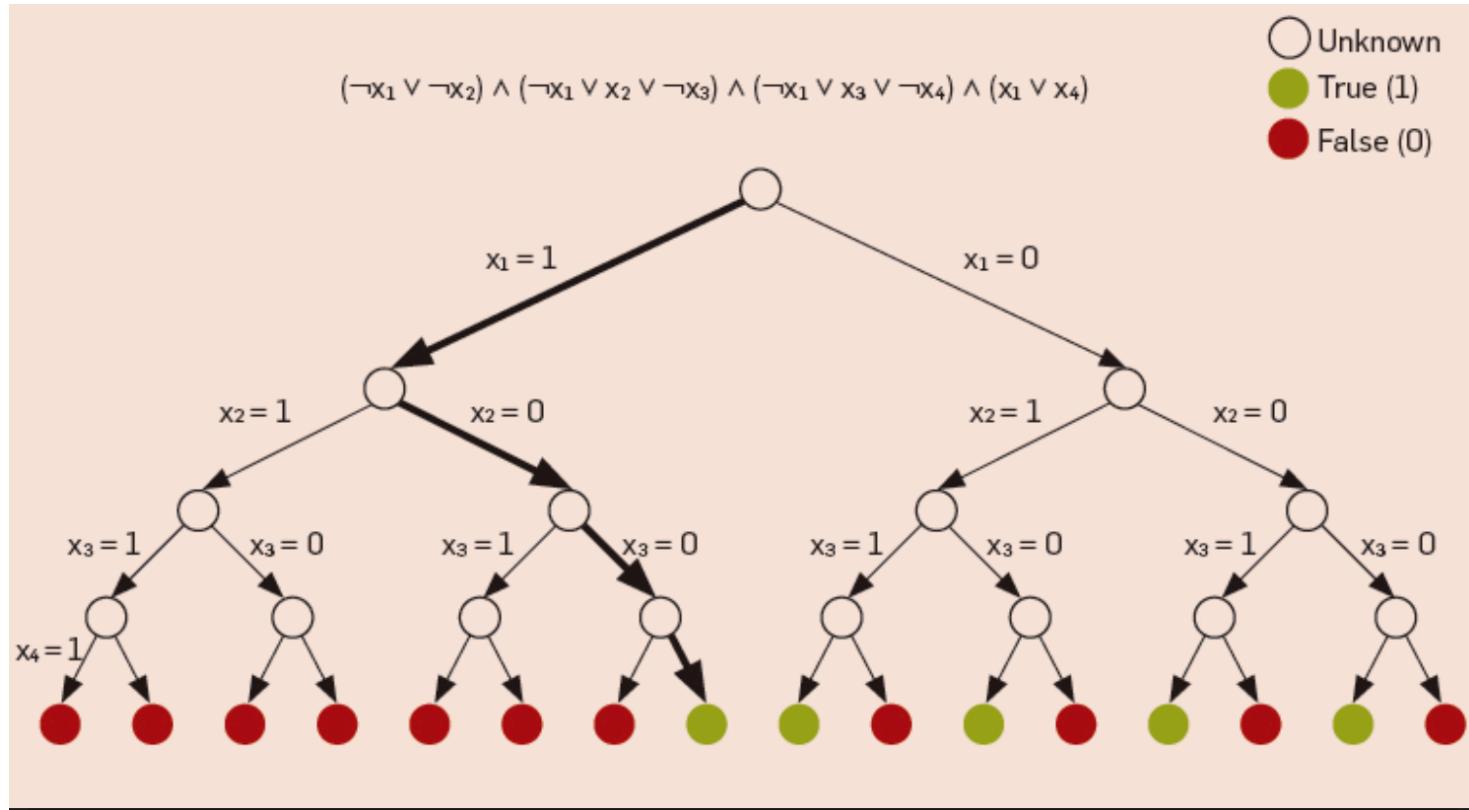


COMPUTATION



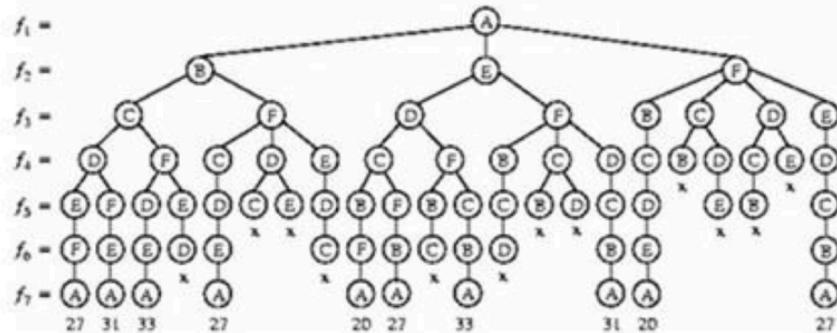
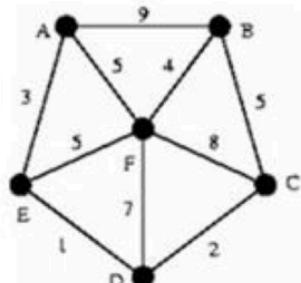
Moore's Law states that the number of transistors that can be placed on an integrated circuit for the same price will increase exponentially by a factor of 2 every 18 to 24 months. In other words, put simply *Moore's Law* claims that CPU processing power will double approximately every two years for the price of 1,000 dollars. (Graph copyright Ray Kurzweil)

NP-HARD and Intractable

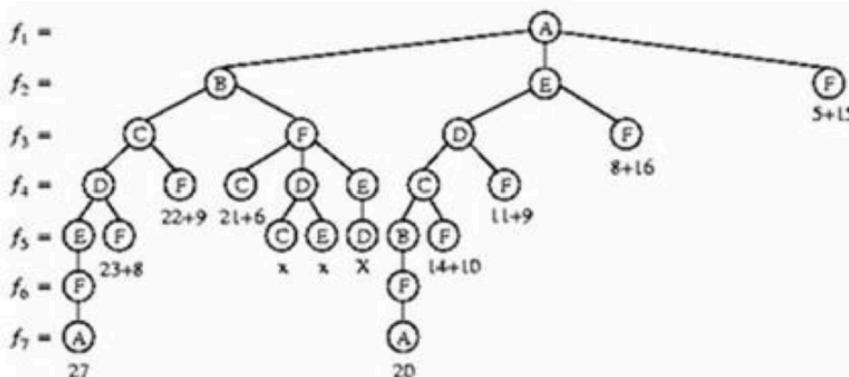


Exhaustive Search v.s. Branch and Bound

TSP example



Backtracking/exhaustive search

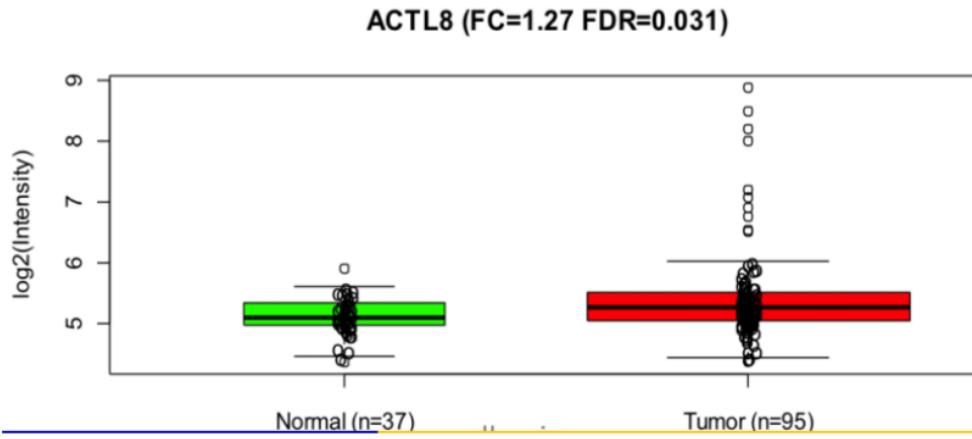


Branch and bound

DATA

At its core, COPA is a technique for analyzing key-value pairs of gene expression data and to detect outliers, which are the candidates for cancer. A parallelized algorithm was necessary owing to the size of the data involved, says Parsian, who also teaches at the University of Santa Clara.

- COPA = Cancer Outlier Profile Analysis
- Statistics designed to identify outliers in cancer gene expression profile

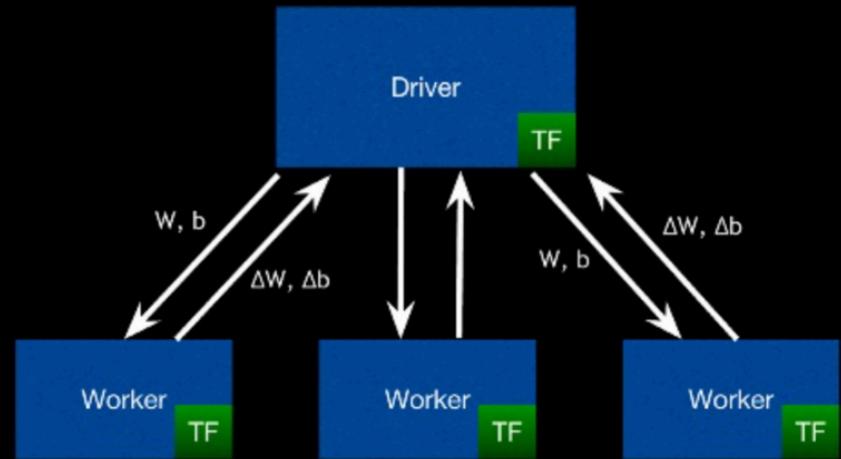
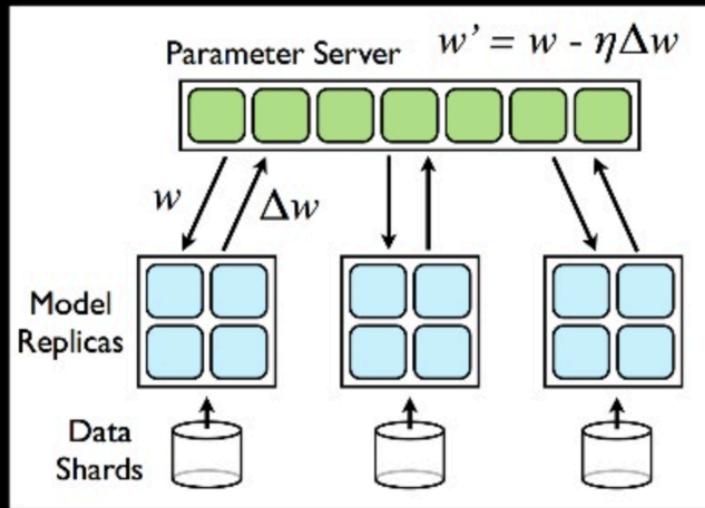


COPA has proven its value in detecting genetic mutations linked to prostate cancer, and is now being studies with other types of cancers

"Of course, if the data size is small, you can detect it visually or by writing some sample programs," Parsian says. "But when you have terabytes of data you're analyzing, detecting mutation is impossible to do visually."

There are many ways to solve this problem, and any number of algorithms could do the

TensorSpark Architecture

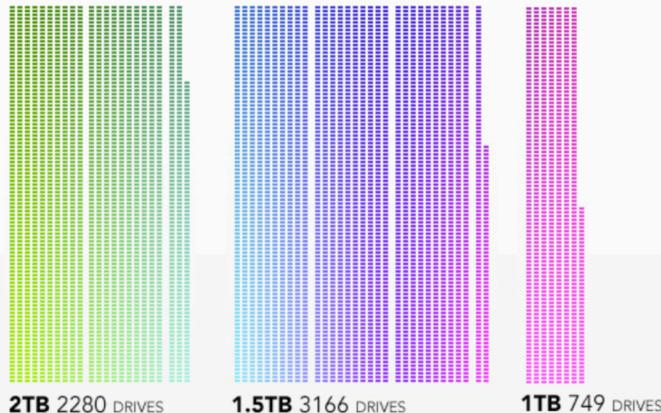


credit: <http://www.slideshare.net/arimoinc/distributed-tensorflow-scaling-googles-deep-learning-library-on-spark>

A PETABYTE IS A LOT OF DATA

- 1 PETABYTE • 20 MILLION FOUR-DRAWER FILING CABINETS FILLED WITH TEXT
- 1.5 PETABYTES • SIZE OF THE 10 BILLION PHOTOS ON FACEBOOK
- 15+ PETABYTES • INTERNET USER'S DATA BACKED UP ON MOZY.COM
- 20 PETABYTES • THE AMOUNT OF DATA PER PROCESSED BY GOOGLE DAY

10 PETABYTES



Dijkstra - Cooperating Sequential Processes

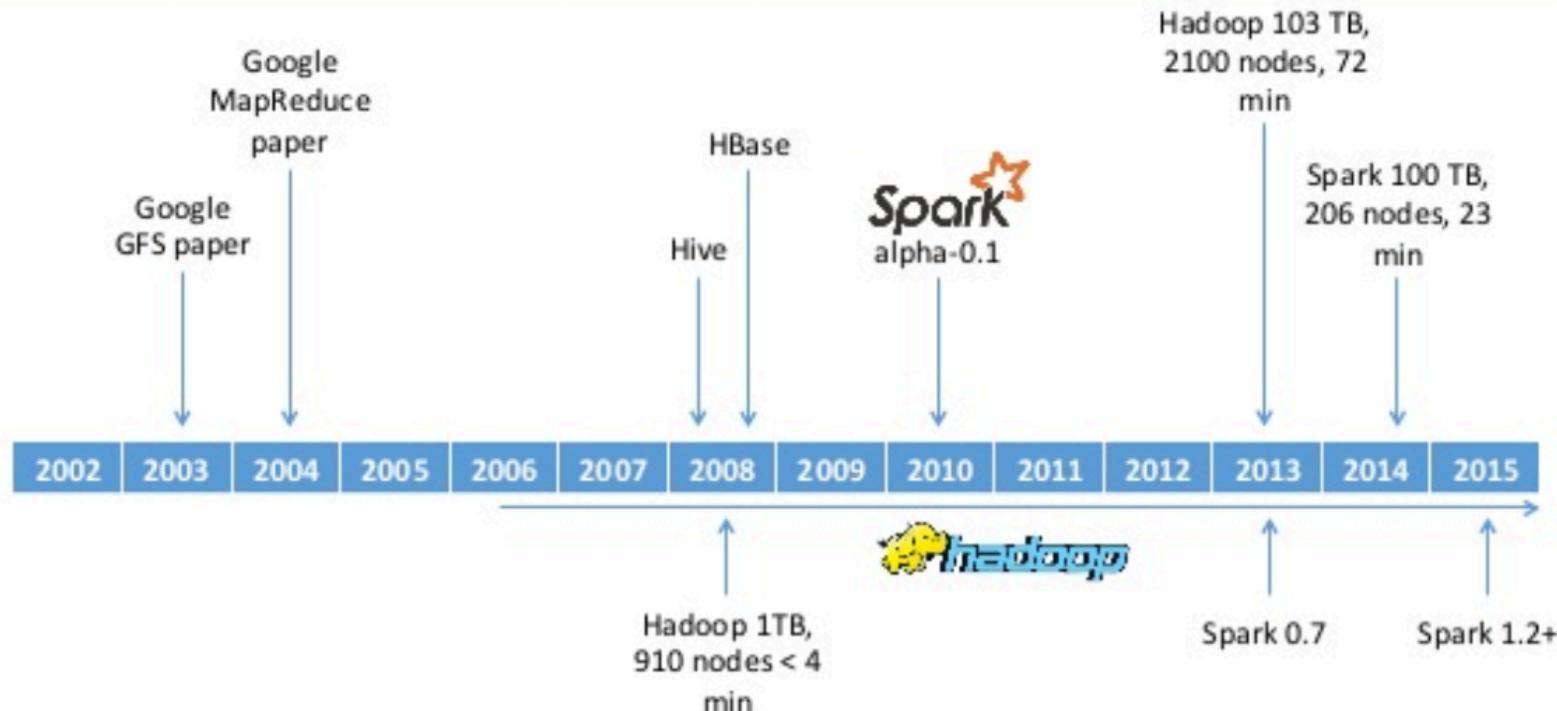


*communicate
cooperate
synchronize*

Caution: Distributed Processing is a hard problem. Keeping track of all the moving parts is challenging.



Timeline



#t3chfest2015

STRATIO

9

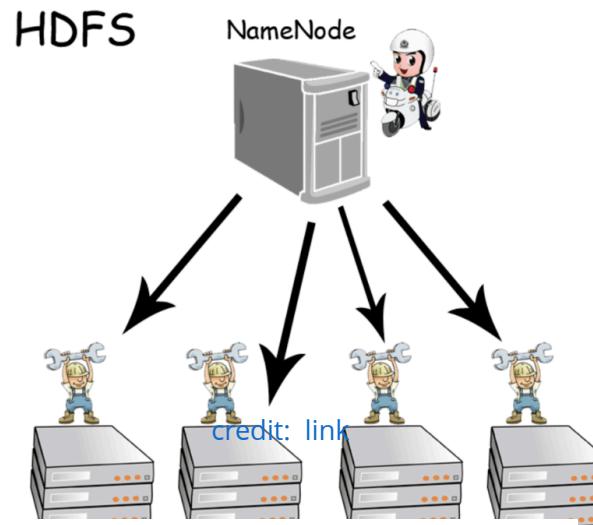
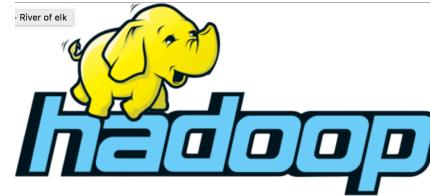
credit : Adios hadoop, Hola Spark! T3chfest 2015

9 of 26



The project includes these modules:

- **Hadoop Common:** The common utilities that support the other Hadoop modules.
- **Hadoop Distributed File System (HDFS™):** A distributed file system that provides high-throughput access to application data.
- **Hadoop YARN:** A framework for job scheduling and cluster resource management.
- **Hadoop MapReduce:** A YARN-based system for parallel processing of large data sets.



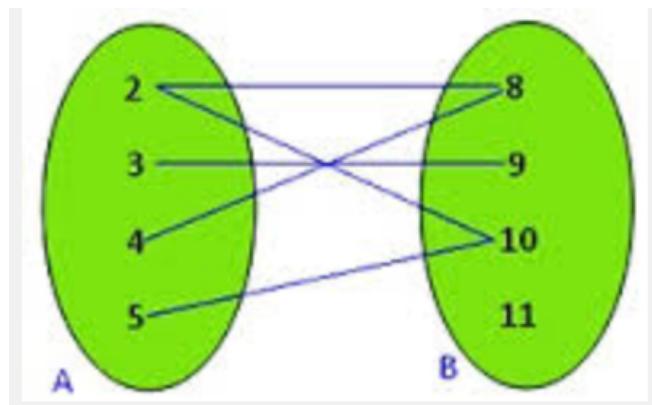


- Began at UC Berkeley in 2009
- Fast and general purpose cluster computing
- 10x faster on disk. 100x faster in-memory
- Integrates with Hadoop and can read existing data
- API's - Java, Python, Scala
- Deeply embraced due to *elegance* of use

Functional Programming

vs

Imperative Programming



Functional



```
range = domain.map(lambda x : x*x)
```

domain is immutable

range is a new fresh object

required for parallel processing

Imperative

```
for i in domain:
```

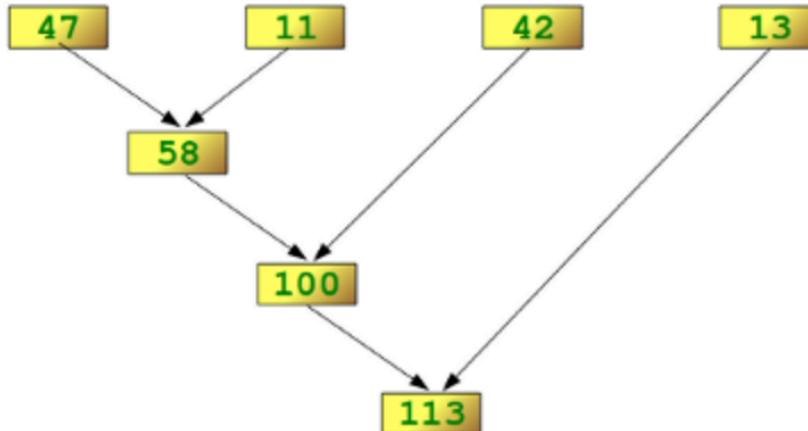
```
    domain[i] = domain[i] * domain[i]
```

danger: side effects, mutations

We illustrate this process in the following example:

```
>>> reduce(lambda x,y: x+y, [47,11,42,13])  
113
```

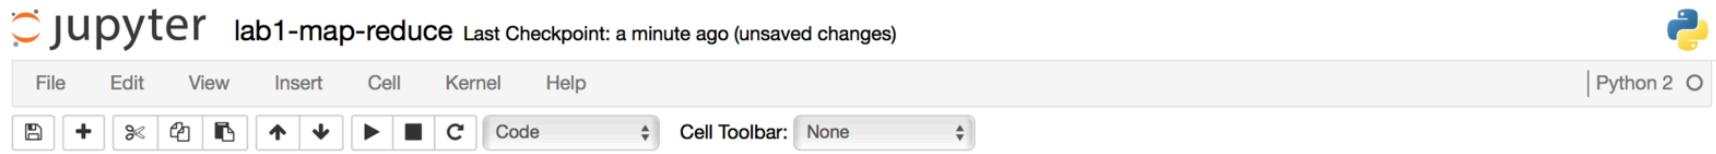
The following diagram shows the intermediate steps of the calculation:



credit: <http://www.python-course.eu/lambda.php>

```
In [17]: reduced = mappedRdd.reduce(lambda x, y: x+y)  
print type(reduced)  
print reduced
```

```
<type 'int'>  
328350
```



Introduction to the Map, Reduce and Filter Abstractions

The Map abstraction

```
In [1]: #Initialize data
nums = range(10)
print (nums)
print (type(nums))
```

```
[0, 1, 2, 3, 4, 5, 6, 7, 8, 9]
<type 'list'>
```

```
In [2]: #Python only
map(lambda x: x*x, nums)
```

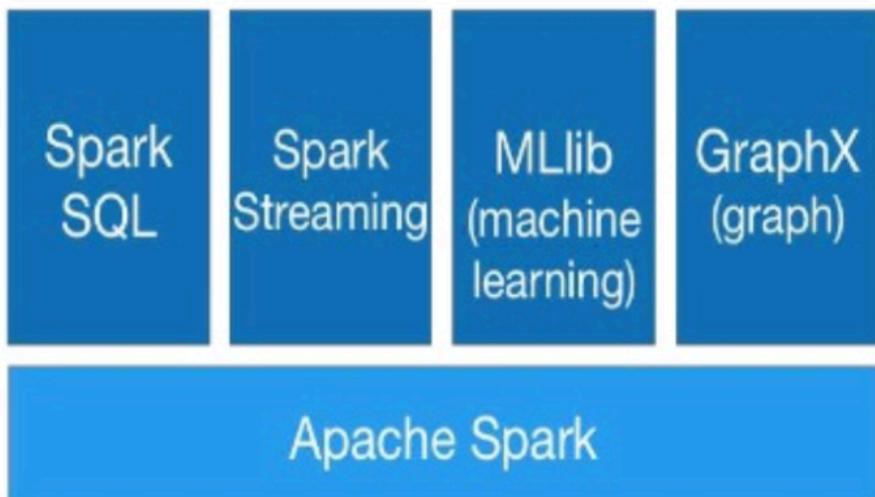
```
Out[2]: [0, 1, 4, 9, 16, 25, 36, 49, 64, 81]
```

```
In [3]: #use a function definition. This is called a closure
def square(x):
    return x*x
```

```
In [4]: #python only
results = map(square, nums)
for num in results:
```

Spark Stack

- Spark SQL
 - For SQL and unstructured data processing
- MLlib
 - Machine Learning Algorithms
- GraphX
 - Graph Processing
- Spark Streaming
 - stream processing of live data streams



<http://spark.apache.org>

1. Most machine learning programs are iterative. Each iteration improves results
2. With MapReduce each iteration is written to disk. This is expensive.
3. Spark runs ***in memory*** using an abstraction known as an **RDD**

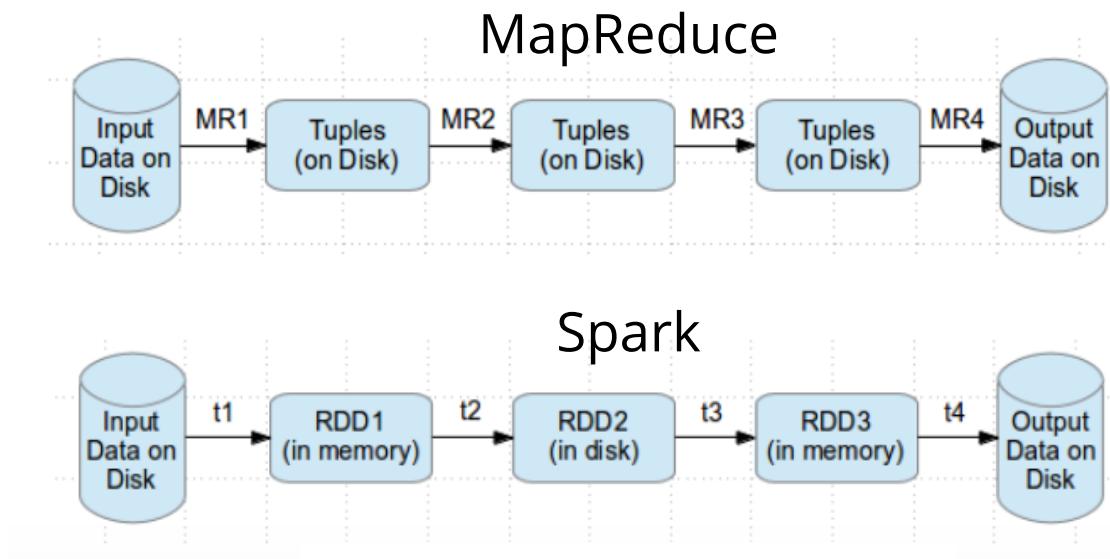
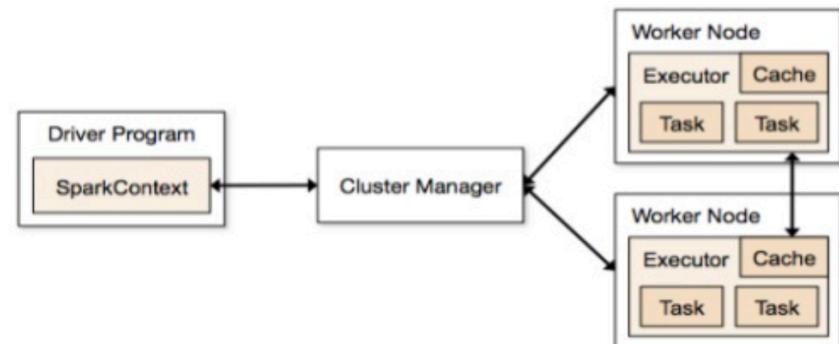
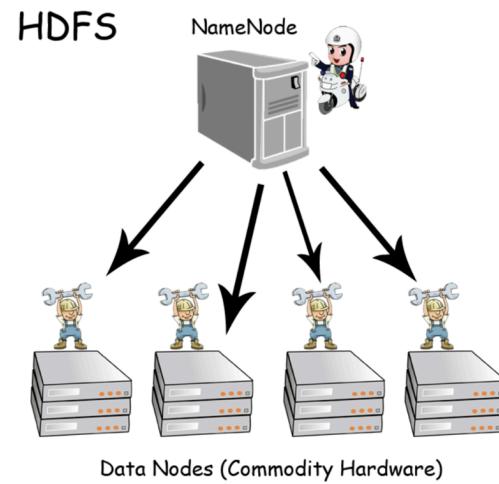


Image Credits: Datatamasha.com

Take the compute to the data

Execution Flow



<http://spark.apache.org/docs/latest/cluster-overview.html>

Coding Exercise: WordCount

```
1 public class WordCount {
2     public static class TokenizerMapper
3         extends Mapper<Object, Text, Text, IntWritable>{
4
5     private final static IntWritable one = new IntWritable(1);
6     private Text word = new Text();
7
8     public void map(Object key, Text value, Context context
9                     ) throws IOException, InterruptedException {
10        StringTokenizer itr = new StringTokenizer(value.toString());
11        while (itr.hasMoreTokens()) {
12            word.set(itr.nextToken());
13            context.write(word, one);
14        }
15    }
16}
17
18 public static class IntSumReducer
19     extends Reducer<Text,IntWritable,Text,IntWritable> {
20     private IntWritable result = new IntWritable();
21
22     public void reduce(Text key, Iterable<IntWritable> values,
23                        Context context
24                        ) throws IOException, InterruptedException {
25
26         int sum = 0;
27         for (IntWritable val : values) {
28             sum += val.get();
29         }
30         result.set(sum);
31         context.write(key, result);
32     }
33}
34
35 public static void main(String[] args) throws Exception {
36     Configuration conf = new Configuration();
37     String[] otherArgs = new GenericOptionsParser(conf, args).getRemainingArgs();
38     if (otherArgs.length < 2) {
39         System.err.println("Usage: wordcount <in> [<in>... <out>]");
40         System.exit(2);
41     }
42     Job job = new Job(conf, "word count");
43     job.setJarByClass(WordCount.class);
44     job.setMapperClass(TokenizerMapper.class);
45     job.setCombinerClass(IntSumReducer.class);
46     job.setReducerClass(IntSumReducer.class);
47     job.setOutputKeyClass(Text.class);
48     job.setOutputValueClass(IntWritable.class);
49     for (int i = 0; i < otherArgs.length - 1; i++) {
50         FileInputFormat.addInputPath(job, new Path(otherArgs[i]));
51     }
52     FileOutputFormat.setOutputPath(job,
53         new Path(otherArgs[otherArgs.length - 1]));
54     System.exit(job.waitForCompletion(true) ? 0 : 1);
55 }
```

```
1 val f = sc.textFile(inputPath)
2 val w = f.flatMap(l => l.split(" ")).map(word => (word, 1)).cache()
3 w.reduceByKey(_ + _).saveAsText(outputPath)
```

WordCount in 3 lines of Spark

WordCount in 50+ lines of Java MR

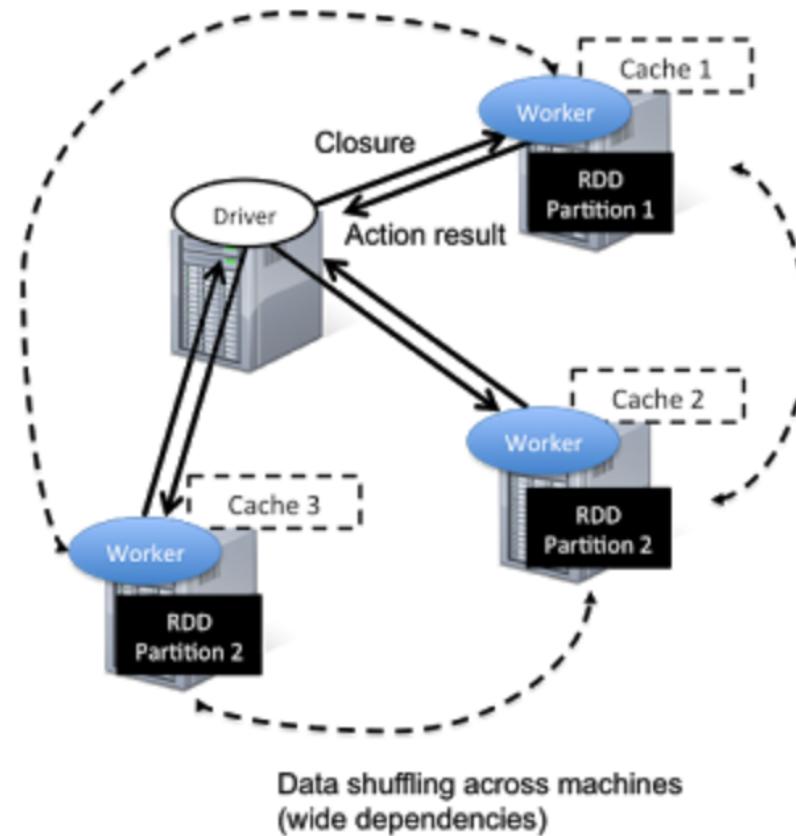
credit: Sparkcamp @ Strata CA: Intro to Apache Spark with Hands-on Tutorials



The Magic of Spark
RDD's
Transformations
Actions

Resilient Distributed Dataset (RDD)

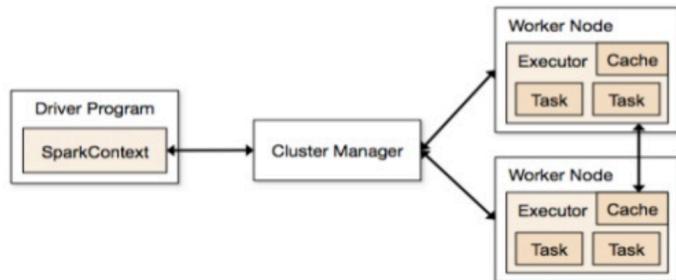
- **RDD** is a basic abstraction
- **Immutable**, partitioned collection of elements that can be operated in parallel
- Basic operations - map, filter, persist
- Multiple implementations - PairRDD <key,value> and Sequence Files



How do I create an RDD?

1. Parallelized Collections
2. External Datasets
3. Streaming Data

Execution Flow



```
data = [1, 2, 3, 4, 5]
distData = sc.parallelize(data)
```

<http://spark.apache.org/docs/latest/cluster-overview.html>

```
>>> distFile = sc.textFile("data.txt")
```

What can I do with an RDD?

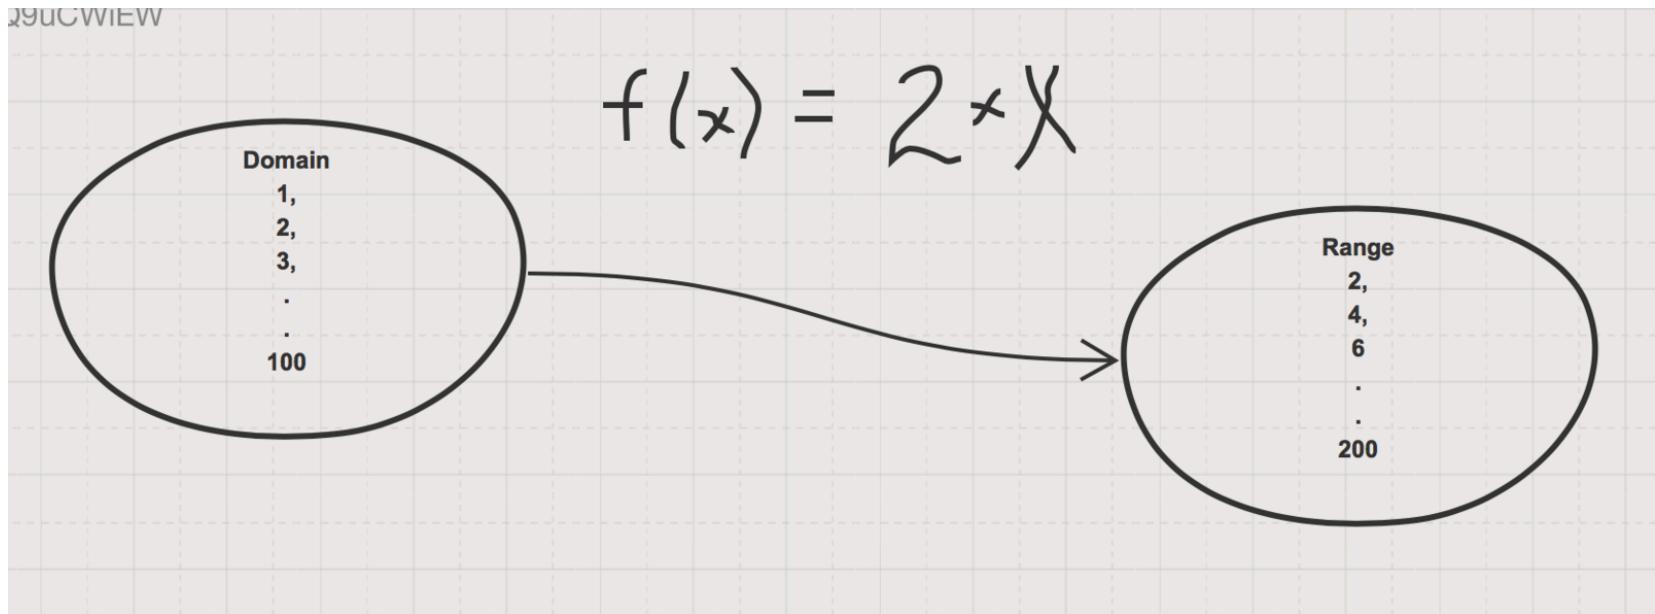
RDD Operations

RDDs support two types of operations: *transformations*, which create a new dataset from an existing one, and *actions*, which return a value to the driver program after running a computation on the dataset.

Map is a transformation that passes each dataset element through a function and returns a new RDD representing the results.

Reduce is an action that aggregates all the elements of the RDD using some function and returns the final result to the driver program.

What functions? Monoids



```
range = domain.map(lambda x : 2*x)
```

```
sum = domain.reduce(lambda x,y : x+y)
```

RDD's are a collection of records

```
rdd = sc.parallelize(range(1000), 5)
```

Transformations create new RDD's from
existing ones

```
errors = rdd.filter(lambda line: "ERROR" in line)
```

Actions materialize a value in the user
program

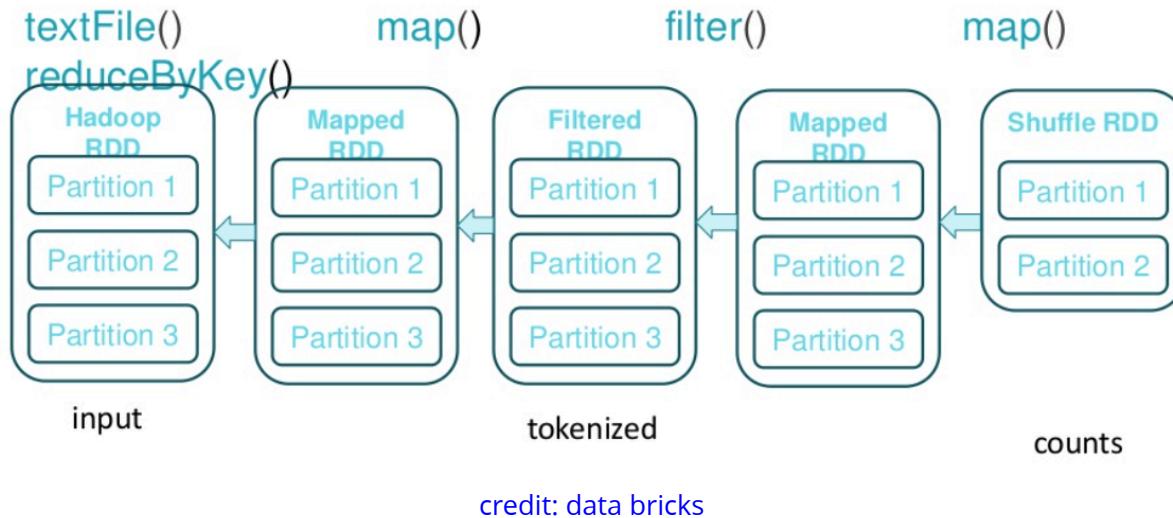
```
size = errors.count()
```

THE DAG

(directed acyclic graph)

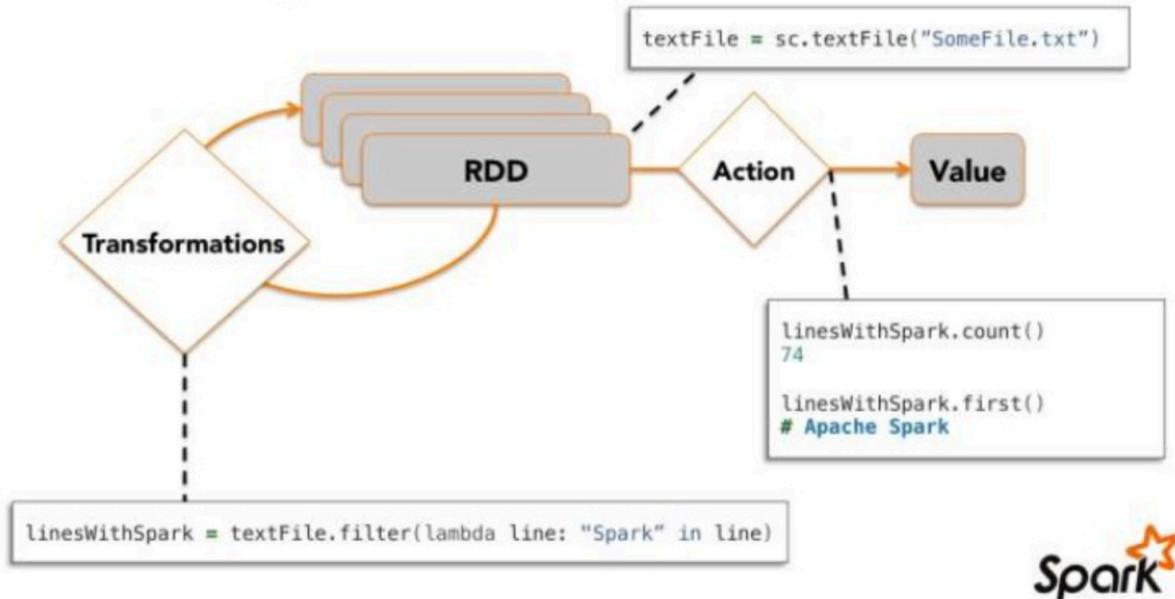
```
sc.textFile().map().filter().map().reduceByKey()
```

DAG View of RDD's



Lazy Evaluation

Working With RDDs



Code runs only upon encountering an action

credit: link

Persistence layers for Spark

Distributed

- Hadoop (HDFS)
- Local file system
- Cassandra
- Amazon S3
- Hive
- Base

File formats

- Text - CSV, Plain Txt
- Sequence File
- AvRO
- Parquet

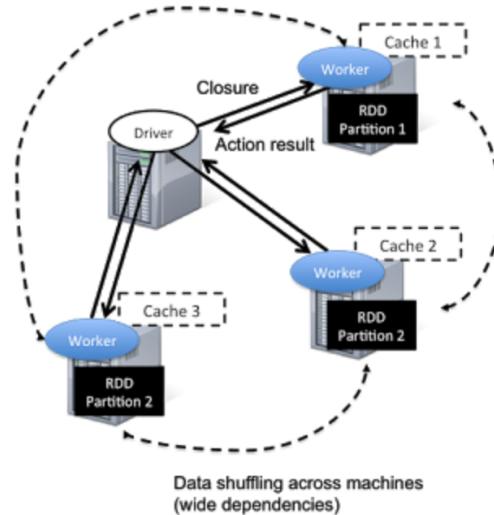
WHEN IS the DAG EXECUTED?

```
import sys
import os

logFile = os.path.join('data', 'logfile')

""" Read and parse log file """
parsed_logs = (sc
    .textFile(logFile)
    .filter(lambda line: 'GET' in line))

print parsed_logs.count()
```



```
in24.inetnebr.com - - 01/Aug/1995:00:00:01 -0400 "GET /shuttle/missions/sts-6.txt HTTP/1.0" 200 1839-
uplherc.upl.com - - 01/Aug/1995:00:00:07 -0400 "GET / HTTP/1.0" 304 0-
uplherc.upl.com - - 01/Aug/1995:00:00:08 -0400 "GET /images/ksc.gif HTTP/1.0" 304 0-
uplherc.upl.com - - 01/Aug/1995:00:00:08 -0400 "GET /images/MOSAIC.gif HTTP/1.0" 304 0-
```

Transformations

<code>map(func)</code>	<code>reduceByKey(func, [numTasks])</code>
<code>filter(func)</code>	<code>aggregateByKey(zeroValue)(seqOp,</code>
<code>flatMap(func)</code>	<code>combOp, [numTasks])</code>
<code>mapPartitions(func)</code>	<code>join(otherDataset, [numTasks])</code>
<code>mapPartitionsWithIndex(func)</code>	<code>cogroup(otherDataset, [numTasks])</code>
<code>union(otherDataset)</code>	<code>cartesian(otherDataset)</code>
<code>intersection(otherDataset)</code>	<code>pipe(command, [envVars])</code>
<code>distinct([numTasks]))</code>	<code>coalesce(numPartitions)</code>
<code>groupByKey([numTasks])</code>	<code>sample(withReplacement, fraction, seed)</code>
<code>sortByKey([ascending], [numTasks])</code>	<code>repartition(numPartitions)</code>

Actions

<code>reduce(func)</code>	<code>take(n)</code>
<code>collect()</code>	<code>takeSample(withReplacement, num, [seed])</code>
<code>count()</code>	<code>takeOrdered(n, [ordering])</code>
<code>first()</code>	<code>saveAsTextFile(path)</code>
<code>countByKey()</code>	<code>saveAsSequenceFile(path)</code>
<code>foreach(func)</code>	<code>saveAsObjectFile(path)</code> (Only Java and Scala)

Key Value Pairs

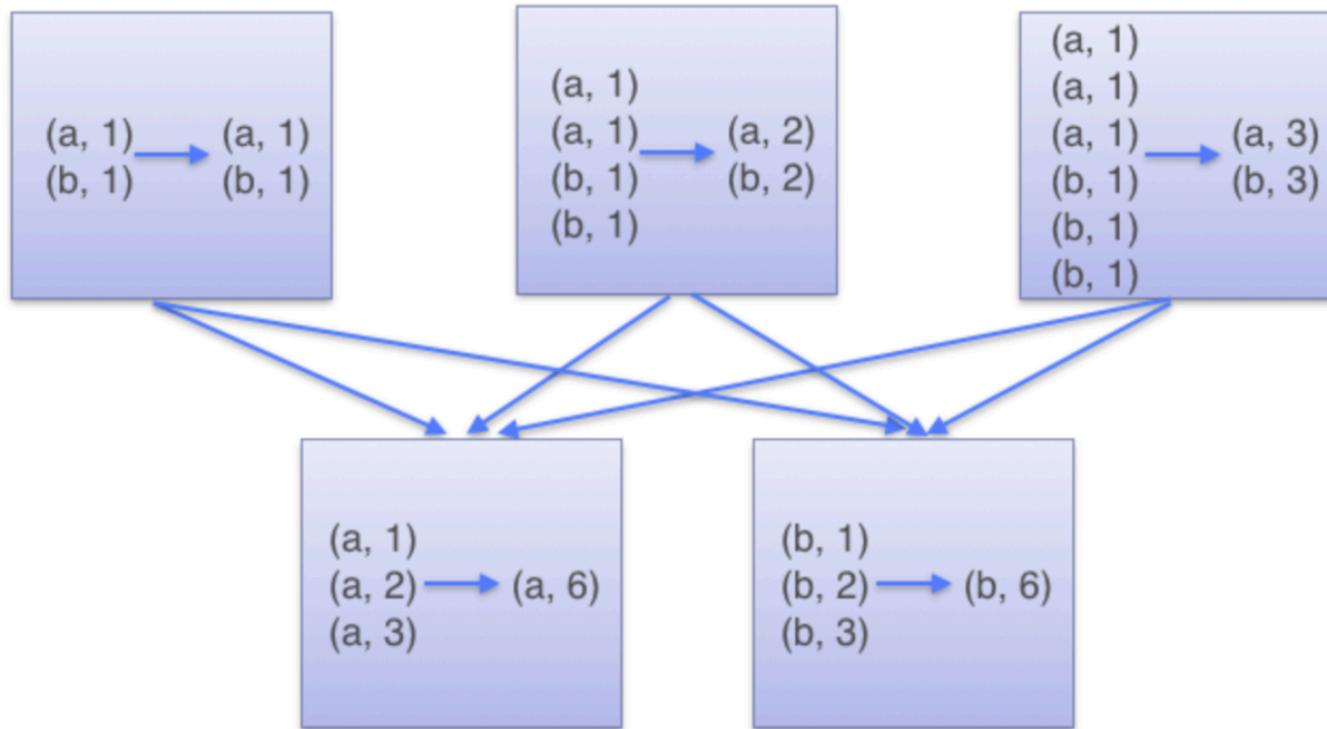
Transformations

- rdd.reduceByKey(func)
- rdd.groupByKey()
- rdd.mapValues(fund)
- rdd.keys()
- rdd.values()
- rdd.sortByKey()

Actions

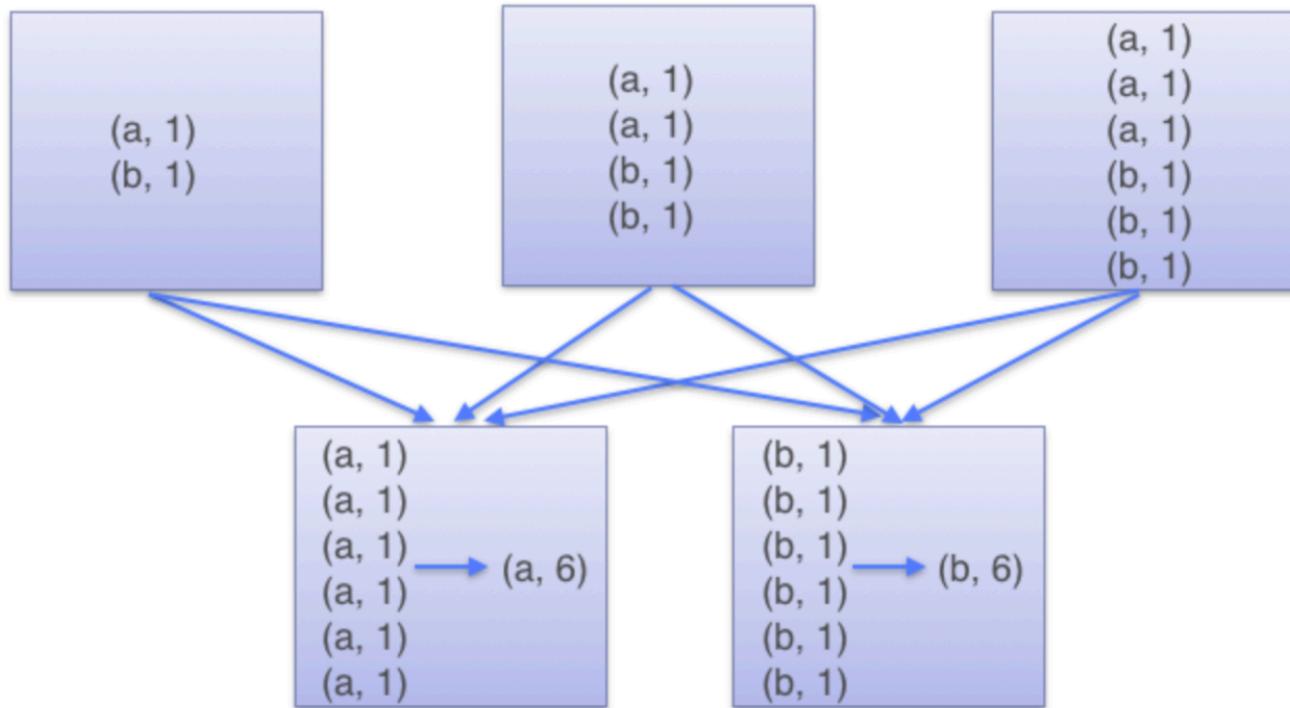
- rdd.countByKey()
- rdd.collectAsMap()
- rdd.lookup(key)

ReduceByKey



credit: https://databricks.gitbooks.io/databricks-spark-knowledge-base/content/best_practices/prefer_reducebykey_over_groupbykey.html

GroupByKey



credit: https://databricks.gitbooks.io/databricks-spark-knowledge-base/content/best_practices/prefer_reducebykey_over_groupbykey.html

```
val lines = sc.textFile("input.txt")
val words = lines.flatMap(line => line.split(" "))
val ones = words.map(s => (s,1))
val count = ones.reduceByKey((a,b) => a + b)
val result = count.collectAsMap()
```

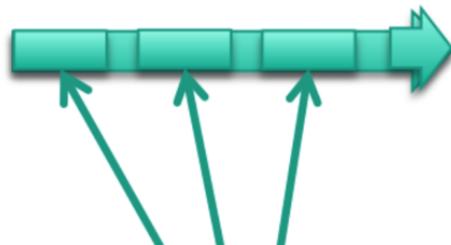
RDD lineage DAG built on driver side
data source RDD
transformation RDD, transformation
action, action RDD

Lab 3 WordCount

Cluster Deployment

- Standalone Deploy Mode
 - simplest way to deploy Spark on a private cluster
- Amazon EC2
 - EC2 scripts are available
 - Very quick launching a new cluster
- Apache Mesos
- Hadoop YARN

live data stream

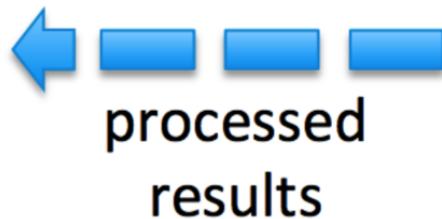


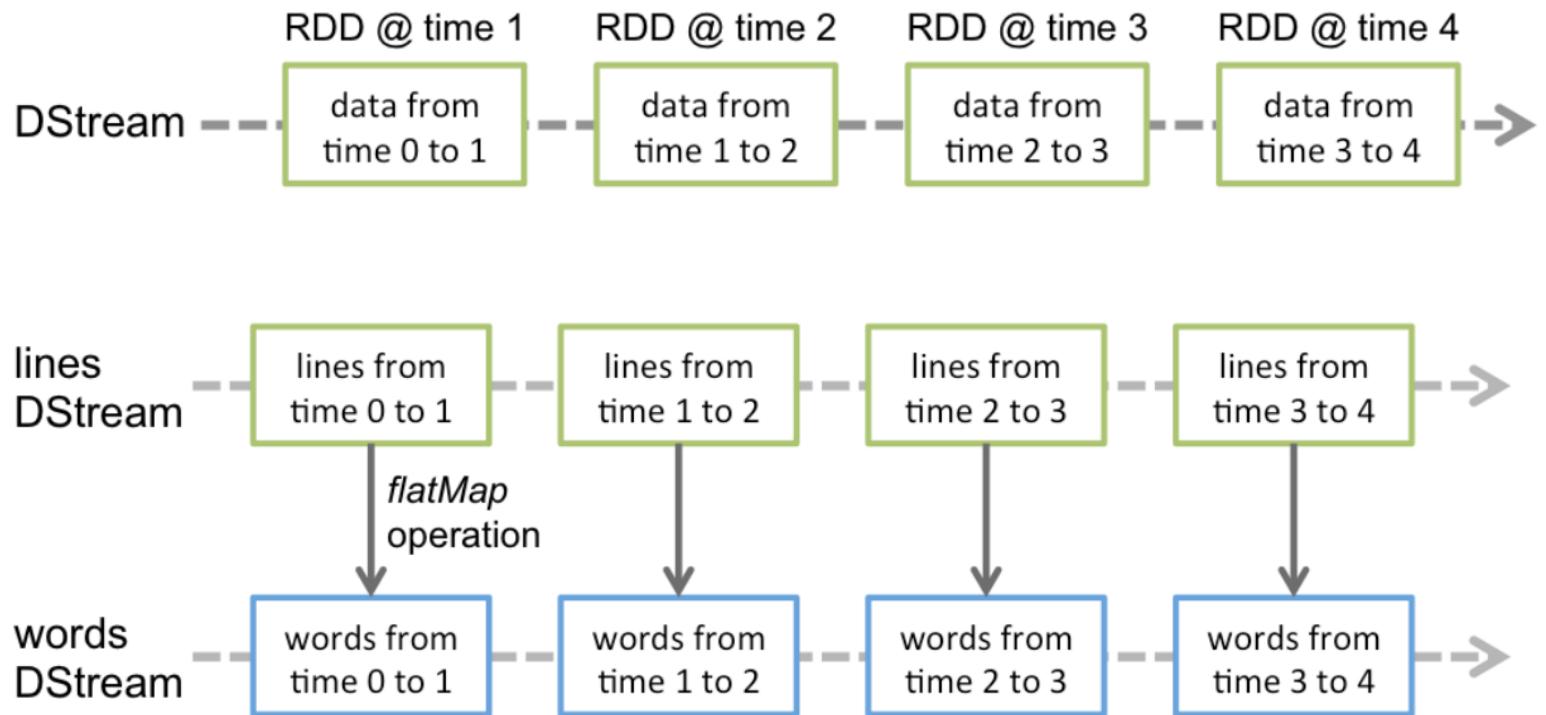
Spark
Streaming

batches of X
seconds

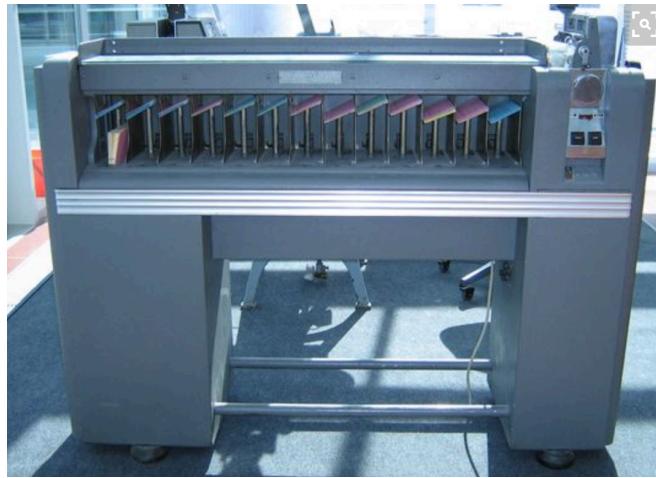
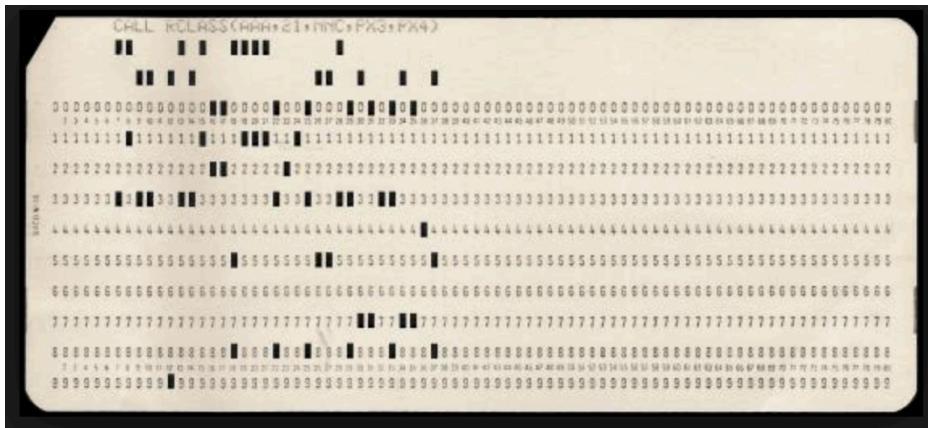


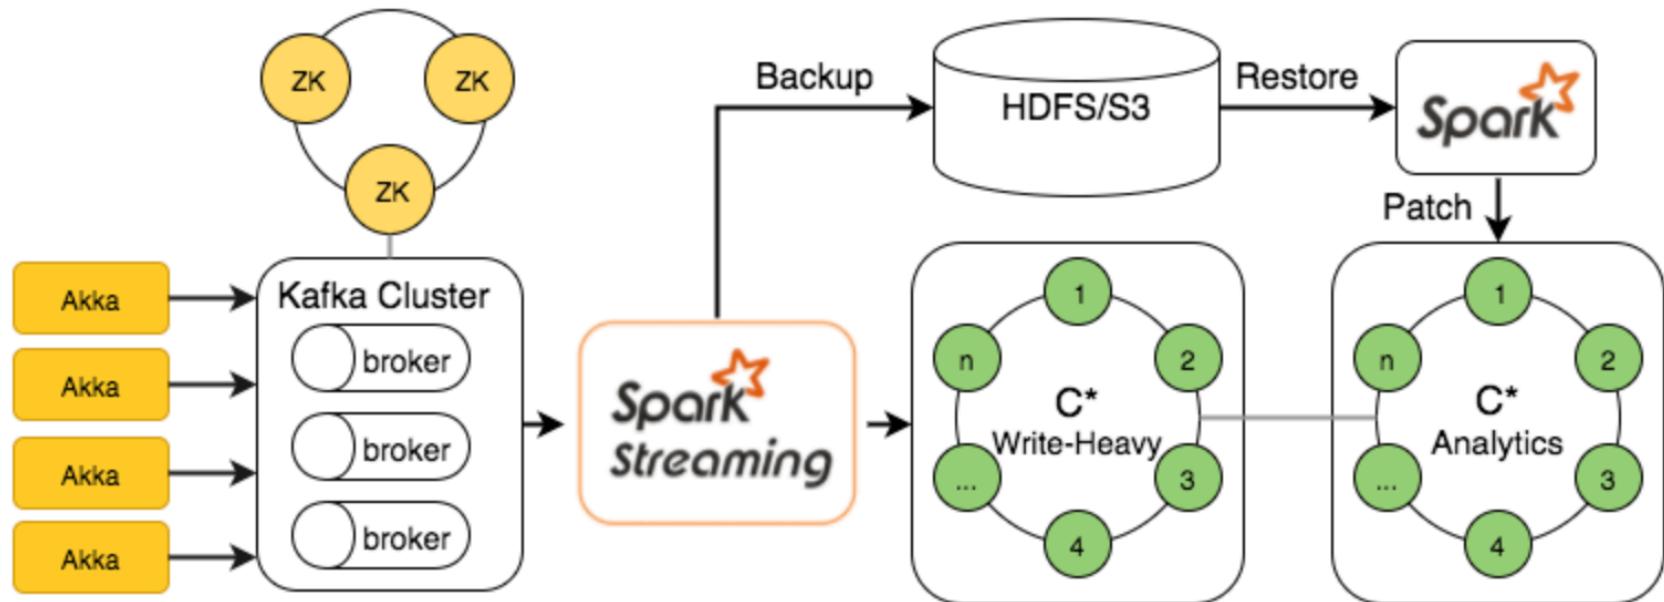
Spark

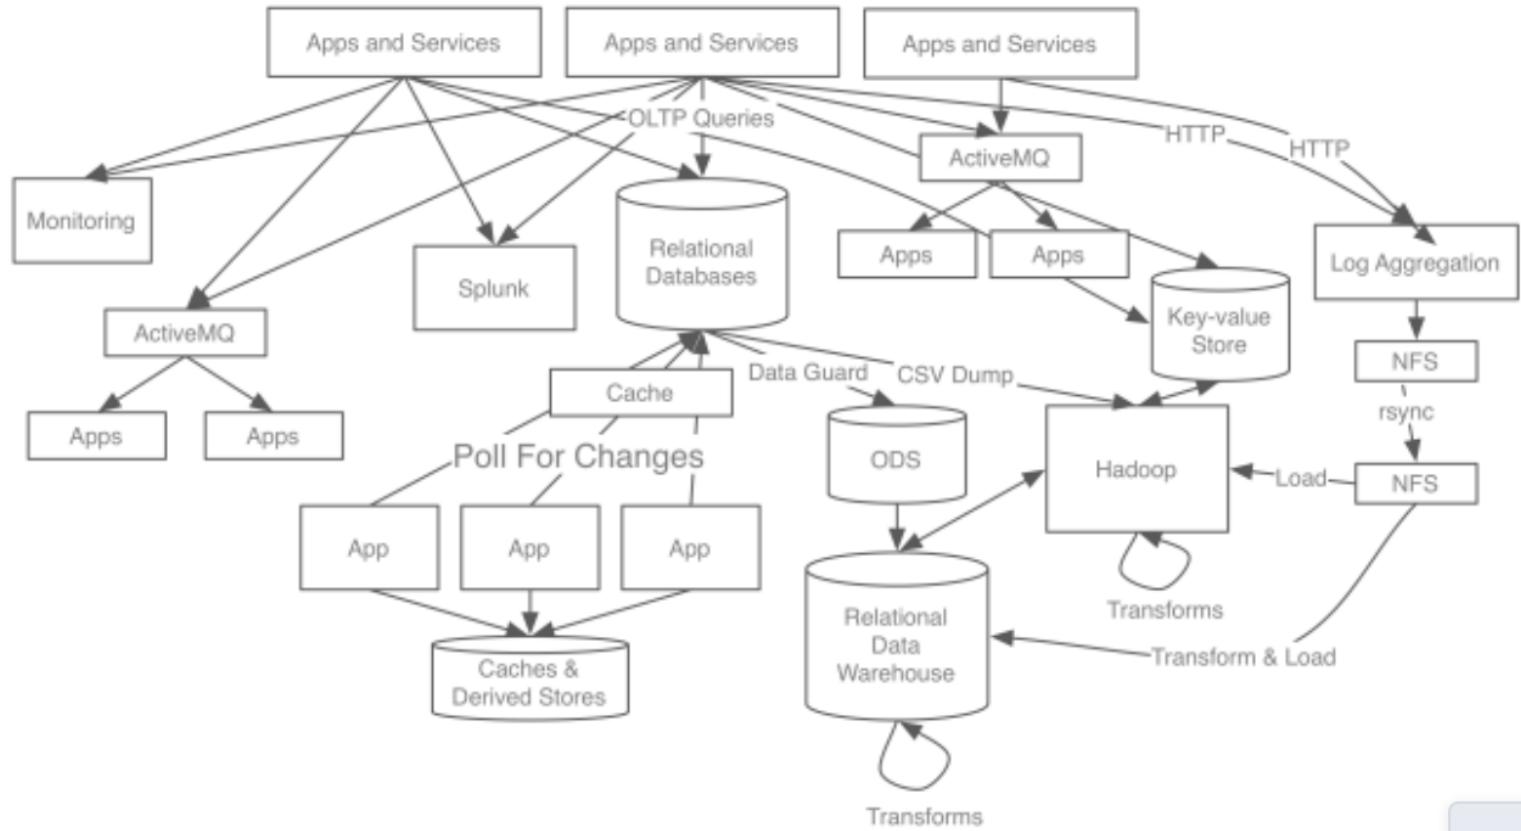


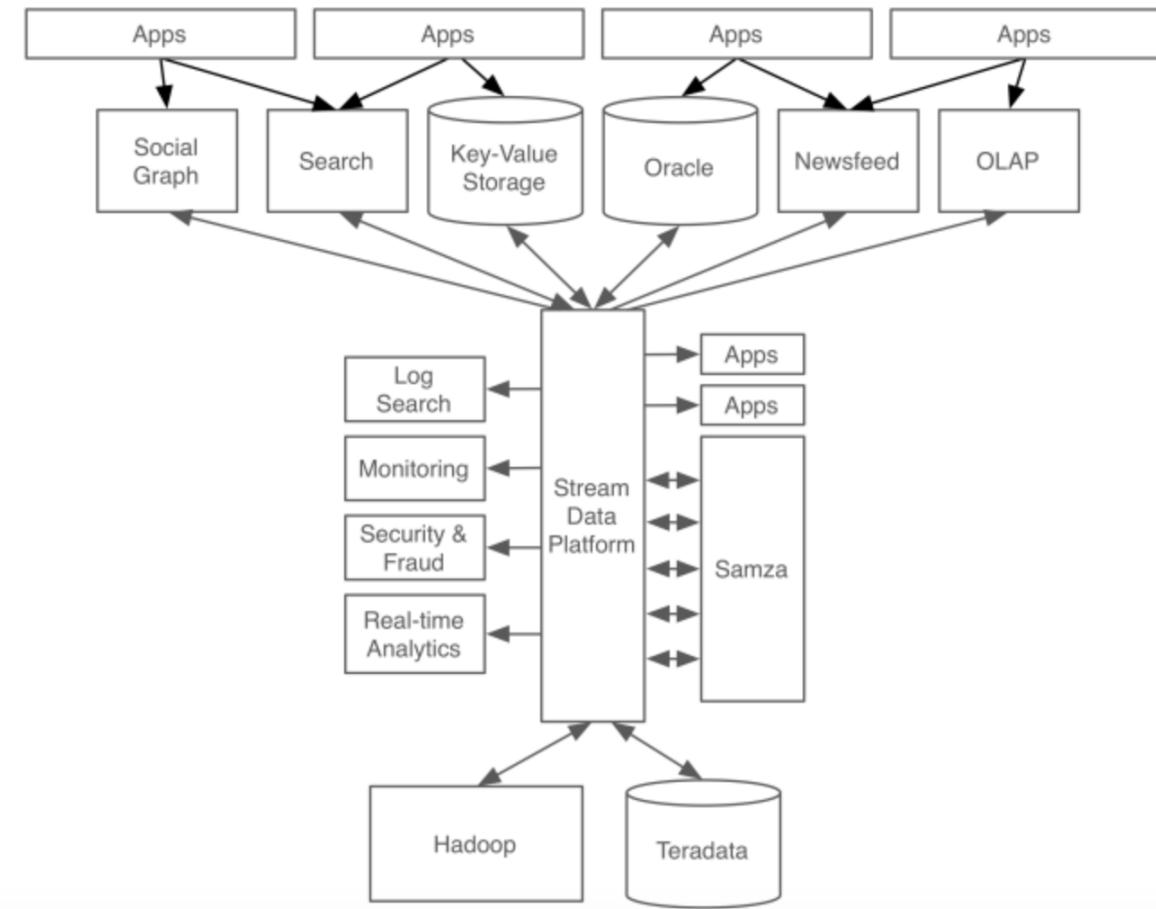












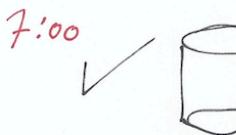


\$500.00

Mary > 10000

6:05 \$ 5.00

6:15 \$ 10.00



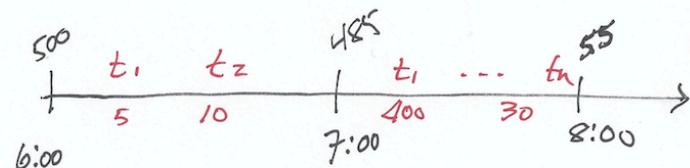
\$ 485.00

7:10 400.00

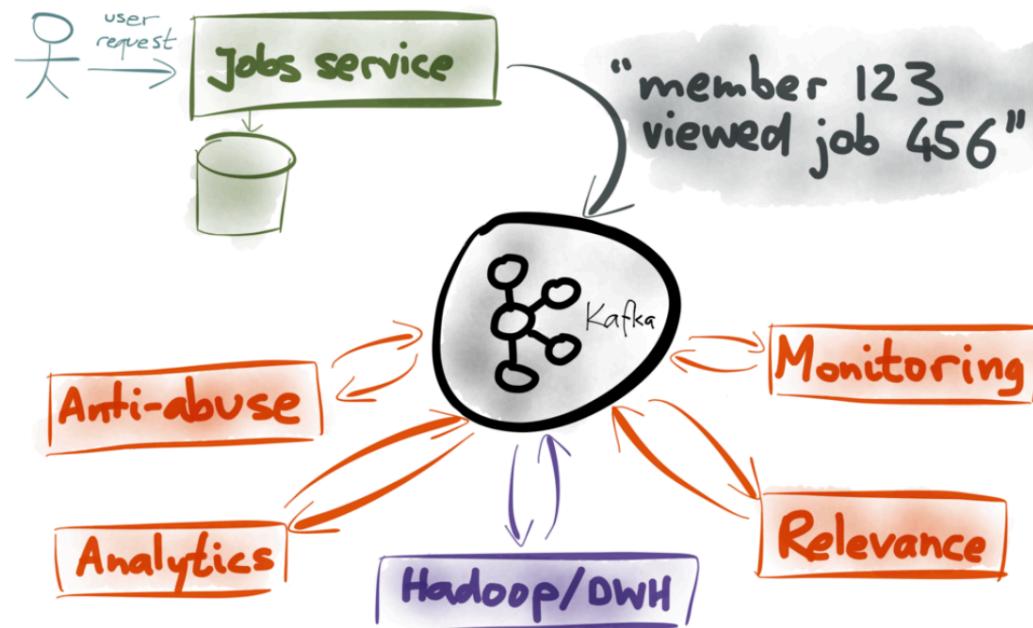
7:30 30.00

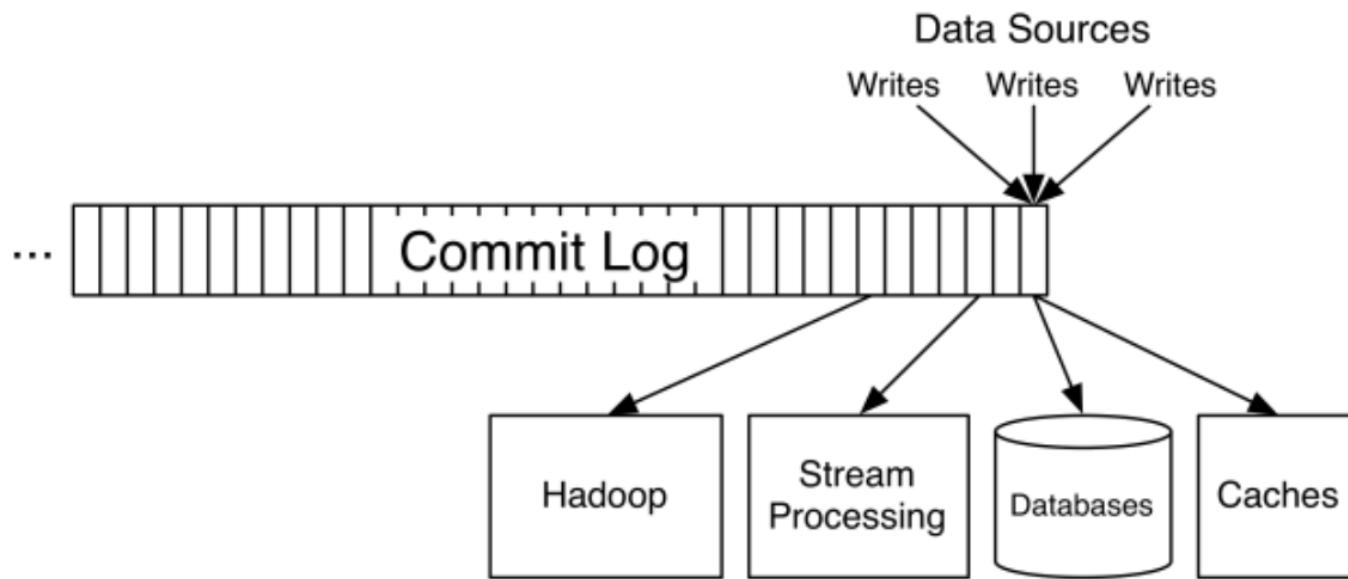


\$ 55.00

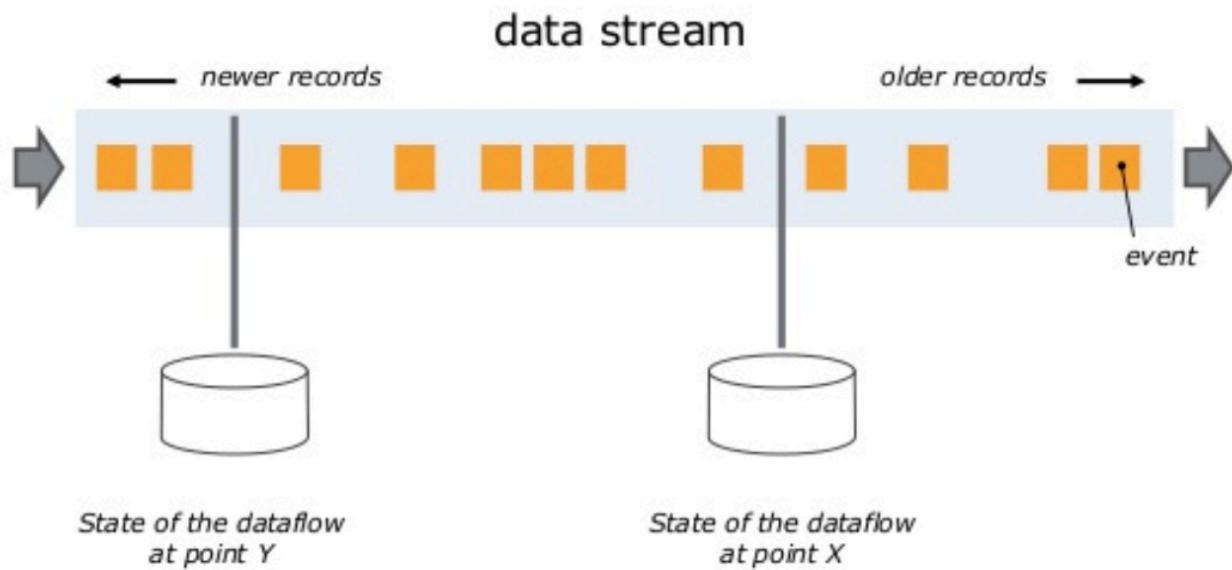


LinkedIn viewing a job posting





Checkpoints



Conclusion

Software is Eating the World